Causal Analysis of Syntactic Agreement Mechanisms in Neural Language Models

Matthew Finlayson

March 15, 2021

Contents

1 Introduction **1**
1.1 Motivation 1

I Background and related work **5**

2 Background and related work **7**
2.1 Targeted Syntactic Evaluation 7
2.2 Probing 8
2.3 Causal Mediation Analysis 9
2.3.1 Total, direct, and indirect effects 9
2.3.2 Causal mediation analysis for neural networks 11

3 Experimental setup **13**
3.1 Data 13
3.2 Models 14

II Experiments and results **17**

4 Experiments and results **19**
4.1 Total Effect 19
4.1.1 Results 21
4.2 Grammaticality Margin 22
4.2.1 Results 23
4.3 Natural Indirect Effect 26
4.3.1 Results 27
4.3.2 Attention Heads 30
4.4 Attention Head Indirect Effects 35

4.5 Adverbs Increase the Probability of Correct Verbs More Than Incorrect Verbs 36
4.6 The (Non-)Impact of Complementizers 36
4.7 Additional Neuron Overlap Details 41
4.7.1 Hypothesizing Syntactic Similarity 41
4.7.2 Neuron Overlap Across Layers 41
4.8 Total Effects Across Architectures 41

5 Conclusion **45**

Abstract

Targeted syntactic evaluations have demonstrated the ability of language models to perform subject-verb agreement given difficult contexts. However, the mechanisms by which neural language models achieve syntactic agreement are not well understood. Thus, this study applies causal mediation analysis to pre-trained neural language models to locate components and discover mechanisms responsible for predicting correctly inflected verbs. In particular, we investigate the magnitude of models' preferences for grammatical inflections, as well as whether neurons process subject-verb agreement similarly across sentences with different syntactic structures. We uncover similarities and differences across architectures and model sizes—notably, that larger models do not necessarily learn stronger preferences. We also observe two distinct mechanisms for producing subject-verb agreement depending on the syntactic structure of the input sentence. Finally, we find that language models rely on similar sets of neurons when given sentences with similar syntactic structure.

Chapter 1

Introduction

This thesis takes a peek inside a number of Transformer language models (LMs) in order to try to better understand the mechanisms that drive their incredible results. The thesis begins with an overview of syntactic trees and number agreement, the task that our language model will be attempting.

To provide a solid foundation for understanding the discussions in the later chapters, this thesis begins with an overview of hierarchical nature of human language. We first discuss

1.1 Motivation

This thesis builds upon a body of existing work interpreting the behavior of neural language models. Targeted evaluations have shown that neural language models (LMs) are able to predict the correct token from a set of minimally differing continuations the vast majority of the time, even in difficult contexts [Linzen et al., 2016, Gulordava et al., 2018]. This is true for a number of linguistic tasks including subject-verb agreement [van Schijndel et al., 2019], filler-gap dependencies [Wilcox et al., 2018], and reflexive anaphora [Marvin and Linzen, 2018]. For example, van Schijndel et al. [2019] show that recurrent neural networks are able to learn subject-verb agreement rules, even when challenged with an attractor, such as in the following example.

(1) The key to the cabinets <u>is</u>/*<u>are</u> next to the coins.

In this case, a model using a linear analysis (i.e., inflecting the verb based on the most recent noun) would favor the ungrammatical inflection, whereas a model using a hierarchical analysis would choose the grammatical inflection.

While we have a reasonable understanding of the generally correct *behavior* of LMs in such contexts, the mechanisms behind the models' sensitivity to syntactic agreement are still not well understood. Recent work has performed causal analyses units within LSTM models. For instance, Lakretz et al. [2019] uncover neurons responsible for syntactic agreement by observing the effects of ablating specific neurons, and Giulianelli et al. [2018] perform causal analyses of LSTM hidden representations' impact on syntactic agreement using "diagnositc classifiers" trained to to predict number from the internal states of the language model. While these studies focus on LSTM models, Transformer-based LMs' agreement mechanisms have not been investigated. Transformer-based LMs' are therefore enticing candidates for further analysis given their superior syntactic generalization abilities compared to those of LSTMs [Hu et al., 2020].

We thus propose applying the behavioral-structural method of causal mediation analysis as put forward by Pearl [2001] to syntactic agreement, following the approach used by Vig et al. [2020] for interpreting gender bias in pre-trained LMs. This method allows us to implicate specific model components in the observed behavior of a model. If we view a neural LM as a causal graph proceeding from inputs to outputs, we can view each model component (e.g., a neuron) as a mediator. We measure the contribution of a mediator to the observed output behavior by performing controlled *interventions* on input sentences and observing how they change the probabilities of continuation pairs. We focus primarily on GPT-2 [Radford et al., 2019], although we also analyze TransformerXL [Dai et al., 2019] and XLNet [Yang et al., 2020].

As a result of our analysis, We find that both GPT-2 and Transformer-XL use two distinct mechanisms to accomplish subject-verb agreement, one of which is active only when the subject and verb are adjacent. Conversely, XLNet shows more unified behavior across syntactic structures. We also find that even though larger models assign a higher probability to the correct inflection more often, this does not necessarily translate to a larger margin between the probability of the correct and incorrect options. Additionally, in larger models, agreement mechanisms are similar to those in smaller models, but more distributed across layers. Finally, we find that the most important neurons for agreement are shared across different structures to various extents, and that the degree of neuron overlap matches well with

human intuitions of syntactic similarity between structures.

Part I

Background and related work

Chapter 2

Background and related work

This chapter gives an overview of previous work that motivates and informs the experiments we perform in part II. We begin with an overview of previous work in syntactic analysis of language models (sections 2.1 and 2.2), then give background for causal mediation analysis in section 2.3.

2.1 Targeted Syntactic Evaluation

As black-box neural language models have rapidly gained popularity and improved over the last decade, interest in evaluating these models' ability to perform non-trivial linguistic tasks has increased. Towards this goal, recent years have seen a number of behavioral studies of neural language models as psycholinguistic subjects [Linzen et al., 2016, Gulordava et al., 2018, Goldberg, 2019], focusing primarily on LSTM LMs [Sundermeyer et al., 2012], GPT-2 [Radford et al., 2019], and BERT [Devlin et al., 2019a][1] as subjects. Among the more popular linguistic phenomena to test are whether models prefer grammatical completions in subject-verb agreement contexts [Marvin and Linzen, 2018, van Schijndel et al., 2019, Goldberg, 2019, Mueller et al., 2020], as well as in filler-gap dependencies [Wilcox et al., 2018, 2019] because these tasks supposedly require a nontrivial knowledge of the hierarchical syntactic structure of language. These targeted syntactic evaluations are largely based on the approach of Linzen et al. [2016], where a model's ability to syntactically generalize is measured by its ability to choose the correct inflection in difficult structural contexts instantiated by tokens that the

[1] BERT is a contextualized word prediction model, not a true language model.

model has not seen together during training. The model's ability to choose the correct inflection is tested by checking whether the model assigns the correct inflection a higher probability than an incorrect inflection given the same context.

The targeted syntactic evaluation approach effectively investigates the output behavior of the model, but does not inform us of inner workings of the model or identify components responsible for the observed behavior. This thesis builds upon this basic approach by incorporating component-level analysis of the model's behavior.

2.2 Probing

A separate line of syntactic analysis work has attempted to look inside black-box neural language models by training classifiers to probe model component outputs. This approach investigates representations within the model associated with syntactic dependencies by defining a family of functions (probes) that map from model representations to some phenomenon that those representations are expected to encode. Hewitt and Manning [2019] and Chi et al. [2020] map BERT representations to dependency parses as a proxy for discovering latent syntactic knowledge within the model. More related, Giulianelli et al. [2018] use probes to investigate how LSTMs handle agreement.

Probing is more difficult to interpret than behavioral approaches because the addition of a trained classifier introduces confounds [Hewitt and Liang, 2019]: most notably, whether the probe maps from model representations to the desired output, or learns the task itself. Probes also only give correlational evidence, rather than causal evidence [Belinkov and Glass, 2019] for component implication in syntactic behavior. Importantly, when using probes, one does not observe behavior in syntactically controlled settings, instead evaluating on tasks such as dependency parsing/labeling; thus, unlike in behavioral studies, we cannot easily make observations about differing model behaviors for different syntactic structures. Our approach, introduced later, attempts to overcome the aforementioned shortcomings by causally implicating model components in syntactic behavior for controlled linguistic structures.

2.3 Causal Mediation Analysis

To accomplish our goals, we turn to the methods of causal mediation analysis. Causal inference methods study the change in a response variable following an intervention; for example, how do health outcomes change after a patient stops consuming nicotine products? Causal mediation analysis [Robins and Greenland, 1992, Pearl, 2001, Robins, 2003] focuses on the role of a mediator in explaining the effect of a treatment on outcomes. For example, if a patient stops using tobacco, are health outcomes mediated by the initial method of nicotine delivery (e.g., smoking tobacco vs. patches vs. nicotine gum)?

This approach lends itself well to interpreting NLP models, as we can view a deep neural network as a graphical model from input to output via mediators, where mediators can be individual components (e.g., neurons). For LMs, the intervention is a change to the input sentence, and the outcome is a function of the probabilities of a set of continuations.

Section 2.3.1 provides an overview of the terms and definitions (an illustration of which can be found in fig. 2.1) used in causal mediation analysis. With this framework, we eventually pursue an interpretation of the mechanisms producing grammatical predictions within neural language models.

2.3.1 Total, direct, and indirect effects

Causal inference methods attempt to infer the causal relationships between variables. These methods have been developed extensively for use in many fields including economics [Maziarz, 2020] and social science [Mahoney, 2010]. The ability to conclusively link the causal effect of one variable on another is particularly valuable in the fields like medicine [Hill, 1965] where the researchers may wish to compare the effects of interventions such as surgery and drugs on patient outcomes. In this section, we use a modified version of the drug treatment scenario from Pearl [2001] as an illustrative example to motivate our own formulation of causal mediation analysis for neural language models. In our example, doctors wish to evaluate the causal effects of a drug treatment intended to treat a disease.

Total effect

In investigating the causal relationship between an independent variable X and dependent variable Y, it makes sense to measure the value of Y as a function of X. This straightforward relationship is known in causal inference literature as *total effect*.

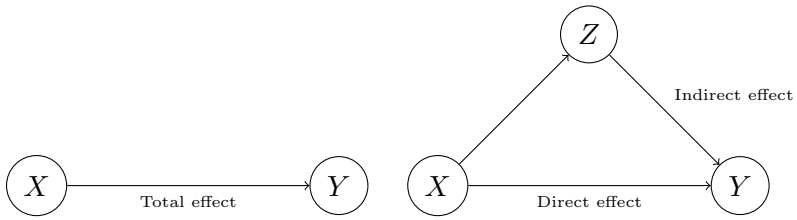

Figure 2.1: Total effect vs. direct effect vs. indirect effect.

As an example, to evaluate the effectiveness of a drug treatment, medical doctors might design a randomized experiment where some proportion of disease patients receive the treatment and the remaining patients receive no treatment. After administering (on not administering) the treatment for a period of time, doctors could measure the effectiveness of the drug by calculating difference in disease rates among those who did and did not receive treatment. If the difference is large, then it can be concluded that the drug is effective. To formalize this, let $X = x$ be the case where the patient receives the treatment, $X = 0$ be the case where the patient receives no treatment, $Y = y_x$ be disease rate for patients receiving the treatment, and $Y = y_0$ be the disease rate for patients receiving no treatment. Doctors may then be interested in the proportional difference

$$\frac{y_x - y_0}{y_0} \tag{2.1}$$

as a metric for the drug's total effect.

Direct effect

Although useful, total effects can be misleading. In our disease treatment example, suppose the drug has a side effect of producing headaches. As a result, patients put on the drug end up consuming more aspirin than those in the control group. If it turns out that aspirin is an effective treatment for the disease, the patients under drug treatment end up with lower disease rates than the control group, despite the possibility the drug itself has no *direct effect* at all. It is therefore important for doctors to be able to identify direct effects of a drug treatment.

In this case, in addition to the independent variable X (treatment) and the dependent variable Y (disease rate), we must contend with a *mediating variable* Z. To find the direct effect of X on Y, controlling for Z, we can

measure the proportional difference in response Y under treatments $X = x$ and $X = x^*$, while holding $Z = z$ constant.

$$\frac{y_{x^*z} - y_x}{y_{0z}} \tag{2.2}$$

where y_{xz} and y_{0z}

In order to find the direct effect of the drug treatment, doctors can control the mediating variable Z by holding aspirin intake constant for both experimental groups and calculate the direct effect as the (proportional) difference in disease rates for patients who receive and do not receive the drug treatment.

Indirect effects

Now that we now can find the direct effect of drug treatment on disease rates, we may now be interested in quantifying the mediating effects of other variables such as aspirin intake. In contrast to total and direct effects, which tell us whether a treatment has an effect, the *indirect* effects of the mediating variables allows us to understand *how* a treatment works.

To find the indirect effect of a variable Z which mediates the effect of independent variable X on Y, we can measure the proportional difference in the response $Y = y$ under the treatment $X = x$ while holding Z to the value z^* it takes when $X = x^*$ and the value z it takes when $X = x$. This gives us

$$\frac{y_{xz^*} - y_x}{y_x} \tag{2.3}$$

In our example, this is the (proportional) difference in disease rates among two groups of patients not given the treatment, where one group takes aspirin at a level that they would have had they been given the treatment, and the other group takes aspirin as they would usually.

2.3.2 Causal mediation analysis for neural networks

In part II, we adapt a methodology developed by Vig et al. [2020] for applying causal mediation analysis (as formulated in Pearl [2001]) to neural networks to study number agreement. In causal mediation analysis for NLP, we view a neural network as a causal graph proceeding from inputs (X), such as linguistic prompts, to output completions (Y), mediated by model components (Z) such as neurons (see fig. 2.2). This setup allows us to implicate specific model components in the behavior of the neural network.

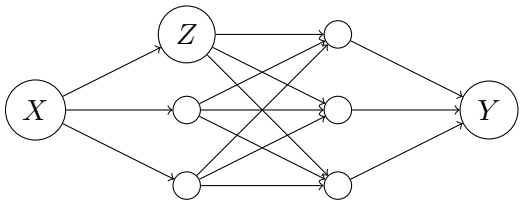

Figure 2.2: A neural network as a causal graph.

In their work, Vig et al. [2020] implicate specific neurons and attention heads in mediating gender bias within various pre-trained LMs. In our work, we do the same for neurons and attention heads mediating syntactic number agreement. This problem is somewhat different because, while one ideally expects equal preferences for male and female completions given gender-ambiguous contexts (for example, given the prompt u "The nurse said that", we want similar continuation probabilities $p(\text{she}|u) \approx p(\text{he}|u)$), this is not the case for subject-verb agreement, where we expect very strong preferences for grammatically correct completions over incorrect completions. A similar assumption of strong preference is also used in recent work by Geiger et al. [2020] who use interventions as a tool to verify that a BERT-based [Devlin et al., 2019b] natural language inference model has learned a specific algorithm.

Chapter 3

Experimental setup

With the goal of discovering which model components are responsible for subject-verb agreement, we begin by constructing a data set of natural language examples where both short- and long-range syntactic dependencies exist between the subject and the verb. We also construct a set of examples where there are 'attracting' nouns that might cause confusion as to which noun the verb ought to agree with. After constructing these examples, we select a number of freely available pre-trained language models on which to perform our interventions.

3.1 Data

As our set of natural language examples, we define a set of prompts $\mathbf{u}$. These prompts are a set of left contexts (beginnings of sentences), generated from a grammar and vocabulary developed by Lakretz et al. [2019]. We opt to synthetically generate prompts rather than sample from a corpus to control for the confound of token collocations in the training set. We modify this grammar to include relative clauses (RCs), generating sentences using the terminals already present. We use prompts from six syntactic structures; an example of each may be found in Figure 3.1. For each structure, we randomly sample 300 prompts. Our dataset, code, and random seeds are available on Github.[1]

In the 'simple agreement' and 'within RC' constructions, there is no separation between the target subject and verb. The 'across one distractor' and 'across

[1] URL withheld for anonymity.

Simple Agreement:
The athlete is/*are

Within Object Relative Clause:
The friend (that) the lawyers *likes/like

Across One Distractor:
The kids gently *admires/admire

Across Two Distractors:
The father openly and deliberately avoids/*avoid

Across Prepositional Phrase:
The mother behind the cars approves/*approve

Across Object Relative Clause:
The farmer (that) the parents love is/*are

Figure 3.1: Syntactic structures used in this study. Ungrammatical forms are marked with asterisks. Target subjects and their agreeing verb inflections are shown in blue, while attractors and their agreeing inflections are shown in red.

two distractors' structures test the effect of placing an adverb/adverbs between the subject and verb. The 'across PP' and 'across RC' structures test the effect of adding a noun (and verb in the latter structure) between the main subject and the main verb. In the 'across RC' and 'within RC' structures, we measure effects both with and without the complementizer *that*.[2]

In each of these constructions, we define a correct and an incorrect continuation. Here, we focus on the third-person singular/plural distinction.

3.2 Models

Our work focuses on a class of neural network language models called Transformers [Vaswani et al., 2017]. A Transformer is a function that takes a sentence (encoded as a vector) as an input, applies a sequence of

[2] A comparison of total and indirect effects when including or excluding the complementizer may be found in Appendix TODO.

linear and nonlinear transformations, and outputs a probability distribution over a vocabulary for the next word (again encoded as a vector). The transformations are parameterized by two sets of model components: neurons and attentions heads. A Transformer is composed of alternating layers of neurons and attention heads, and, in general, the neurons greatly outnumber the attention heads, with hundreds of neurons and only a few attention heads per layer. The exact details of how these components work are beyond the scope of this thesis, but it suffices to understand that "weights", or settings of the neurons and attention heads determine the behavior of the Transformer. Machine learning algorithms use massive amounts of training data to fine tune the neuron and attention head parameters to achieve state-of-the-art performance on language modeling tasks.

In this thesis, we focus primarily on GPT-2 [Radford et al., 2019], an autoregressive Transformer-based English LM. We use several sizes: small, medium, large, and extra-large. We also use DistilGPT-2 [Sanh et al., 2020], a very small distilled version.

To investigate how differences in training across Transformer-based architectures manifest themselves in syntactic agreement mechanisms, we also investigate Transformer-XL [Dai et al., 2019] and XLNet [Yang et al., 2020]. Transformer-XL is an autoregressive English LM whose training objective is similar conceptually to GPT-2's; however, it has a much longer effective context due to its segment-level recurrence. XLNet is an English LM which proceeds through all permutations of the input tokens during training, and which uses a distinct attention masking mechanism as well; during testing, it proceeds autoregressively through the input similar to the other two models.

Part II

Experiments and results

Chapter 4

Experiments and results

Using methods of causal mediation, we perform a number of experiments to investigate a number of effects on various components of the model. We begin with total effect, calculating the difference in probabilities for different verb inflections for prompts of different number. To further understand these effects, we further decompose the total effect to discover the source of the differences across syntactic structures. We then investigate the natural indirect effect, discovering, among other things, a bifurcated mechanism within some Transformer models for performing number agreement.

4.1 Total Effect

In causal analysis, the total effect is defined as the effect of changing some independent variable X on some dependent variable Y. In our experiments, we are interested in how strongly models prefer correct verb forms. To measure this, we first identify modifications to the prompt fed to the model as the independent variable, and the output probabilities of a verb's number inflections as the dependent variable. We can then obtain the total effect by comparing the model prediction probabilities of verbs when preceded by a singular subject to the probabilities when preceded by a plural subject.

To formalize this, we need to define a measure for a model's preference for correct verb forms. Given a verb $v \in \mathbf{v}$ and prompt $u \in \mathbf{u}$ with number feature *sg*, we use the relative prediction probabilities of the grammatically correct v_{sg} and incorrect v_{pl} verb forms as a measure of the preference of a

model with parameters θ for the grammatically correct inflection:

$$y(u, v) = \frac{p_\theta(v_{pl} \mid u_{sg})}{p_\theta(v_{sg} \mid u_{sg})} \tag{4.1}$$

where $y < 1$ indicates a preference for the correct inflection given a singular prompt, and $y > 1$ indicates a preference for the incorrect inflection.[1]

To obtain our independent variable, we define a class of interventions $\mathbf{x}$ that modify the prompts in $\mathbf{u}$ in a systematic way. As we are eventually concerned with the ability of models to choose correct inflections despite the presence of distractors and attractors, we define the intervention `swap-number`, which replaces the prompt subject with the same lexeme of the opposite number inflection (e.g., change "author" to "authors" or vice versa) u_{pl}. We also define the `null` intervention, which leaves u_{sg} as-is (following the methods of Vig et al. 2020).

Returning to our dependent variable, we now define $y_x(u, v)$ as the value of y under intervention x on prompt u. Because the intervention `swap-number` entails swapping the subject for a noun of the opposite number, we now expect $y_{\texttt{swap-number}} > 1$ in Equation 4.1 if the model behaves as expected, since the verb that was originally the correct inflection is now incorrect and vice versa.

Having established out independent and dependent variables, we now define the total effect (TE) for prompt u as the relative difference between the probability ratio y under the `swap-number` intervention and the ratio under the `null` intervention:

$$\begin{aligned}
\mathrm{TE}(\texttt{swap-number, null}; y, u, v) &= \frac{y_{\texttt{swap-number}}(u, v) - y_{\texttt{null}}(u, v)}{y_{\texttt{null}}(u, v)} \\
&= \frac{y_{\texttt{swap-number}}(u, v)}{y_{\texttt{null}}(u, v)} - 1 \\
&= \frac{p_\theta(v_{pl} \mid u_{pl})}{p_\theta(v_{sg} \mid u_{pl})} \Big/ \frac{p_\theta(v_{pl} \mid u_{sg})}{p_\theta(v_{sg} \mid u_{sg})} - 1
\end{aligned} \tag{4.2}$$

We interpret this quantity as the overall preference of a model for the correct inflection of v in context u. Since we observe both singular and plural versions of the prompt, it is arbitrary which grammatical number we start with.

[1] We arbitrarily choose to start with *sg*; we can swap *sg* and *pl* in Eq. 4.1 without loss of generality since we do not directly observe y. This is clarified after Eq. 4.2.

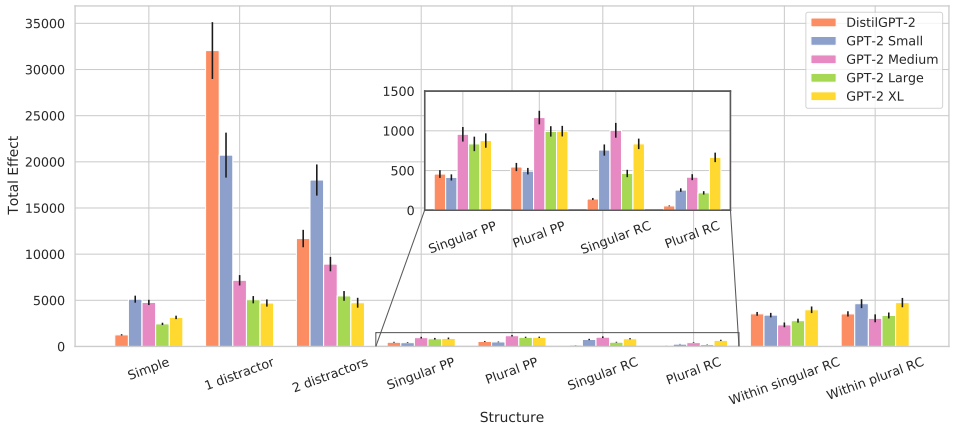

Figure 4.1: Total effects for each structure by model size for GPT-2. Adverbial distractors increase total effects, while attractor phrases decrease them.

In our analysis, we are interested in the average total effect across prompts and verbs:

$$\overline{\mathrm{TE}}(\texttt{swap-number, null}; y) = \mathbb{E}_{u,v}\left[\frac{y_{\texttt{swap-number}}(u,v)}{y_{\texttt{null}}(u,v)} - 1\right] \tag{4.3}$$

We calculate the average total effect for each syntactic construction for different sizes of GPT-2. Unlike in Linzen et al. [2016], we do not measure accuracies by checking whether one probability is higher than another. Rather, the total effect quantifies the margin between the probabilities of correct and incorrect continuations with some intervention.

Because larger models tend to exhibit correct subject-verb agreement more often than smaller models [Hu et al., 2020, van Schijndel et al., 2019], we hypothesize that larger models will generally have larger TEs for the same structure (i.e., we predict that higher accuracies are indicative of larger margins between probabilities).

4.1.1 Results

Figure 4.1 presents total effects by structure for various sizes of GPT-2. With randomized weights, TEs are always near-zero, as expected, and as such are not shown in the figure.

In 'simple agreement' and 'within RC' (where there is no separation of subject and verb, TEs vary between 1,000 and 5,000, depending on model size. This is far higher than the TEs below 250 reported for gender bias in Vig et al. [2020], which is to be expected: GPT-2's training objective explicitly optimizes for predicting grammatically correct tokens given context, but does not optimize for unbiased gender assignment.

Analyzing our results, we arrive at a number of conclusions.

Compared to small models, larger model do not have a stronger preference for correct inflections. Contrary to our hypothesis, and unlike Vig et al. [2020], we do not observe larger TEs for larger models, indicating that models do not necessarily prefer correct inflections more as models scale. While larger GPT-2 models have been shown to achieve higher accuracy on targeted syntactic evaluation sets [Hu et al., 2020], our results show that this performance advantage is not necessarily due to a larger preference for correct inflections.

Adverbial distractors increase correct inflection preference. TEs are highest for structures where distractors are present, with DistilGPT-2 and GPT-2 Small attaining the highest TEs in such contexts. This is surprising, as we expect subject-verb agreement accuracy to decline as the distance between the subject and the verb increases. We suspect the adverbs are acting as cues that a verb will soon appear, thus increasing the probability of both the correct and incorrect verb, but increasing that of the correct verb more. Additional analysis confirms that verb probabilities are notably higher in the distractor structures; see TODO

Interposing attractors decrease correct inflection preference. When PPs or RCs separate the subject and verb, TEs decrease. The number of the attractor does not significantly change TEs across PPs, but does have a more notable effect across RCs: GPT-2 is more certain of its choices across singular RCs than across plural RCs, as evidenced by higher TEs for the former. This may point to a bias within the model, which will be discussed further in section 4.2.

4.2 Grammaticality Margin

Noticing that the plurality of attractors makes a difference in total effects raises the question of whether models have equal preference for correct

inflections when given singular and plural subjects. In other words, whether singular and singular agreement are treated equally by language models.

Because total effects do not disentangle model preferences attained when the subject is singular or plural, we also propose the grammaticality margin metric (shortened to grammaticality hereafter) to aid in the total effect analysis. This metric allows us to observe whether inflectional preferences are larger for singular or plural prompts, whereas total effects measure preferences without respect to grammatical number. Grammaticality is defined as the probability ratio of the correct inflection over the incorrect inflection given u with a specific number feature sg or pl:

$$G(u_{sg}, v) = \frac{p_\theta(v_{sg}|u_{sg})}{p_\theta(v_{pl}|u_{sg})} \tag{4.4}$$

This is equivalent to the reciprocal of $y_{\texttt{null}}$ in the u_{sg} case. Unlike total effects, whether u starts as singular or plural *does* matter here, since we are explicitly measuring the effect of grammatical number on inflectional preferences.

We expect grammaticality values to be higher when the subject and attractors match in number than when they do not match. We also expect grammaticality values to be similar for singular and plural subjects; however, differences can indicate biases toward a certain grammatical number in general.

4.2.1 Results

Throughout this section, we often refer to structures as "having" grammaticality values for convenience. A more accurate description would say that the output of the model has some grammaticality value when the input has a specific syntactic structure.

Plural subjects yield higher grammaticality. Figure 4.2 presents grammaticality values separately for singular and plural subjects, as well as singular and plural attractors when applicable. It is immediately apparent that plural subjects always have higher grammaticality values regardless of the structure or attractor number. In other words, GPT-2 forms stronger agreement dependencies between verbs and *plural* subjects than singular subjects.

We speculate that this phenomenon may be due to the nature of the number inflection in English verbs. In English, the present-tense third-person plural

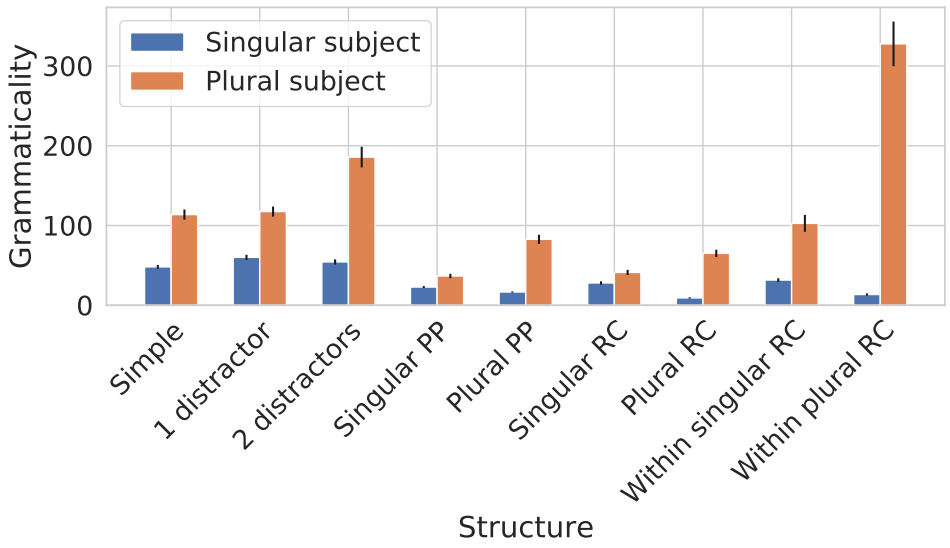

Figure 4.2: Grammaticality for each structure for GPT-2 Medium. The subject number (indicated by bar color) refers to the grammatical number of the subject with which the target verb agrees; the number in the structure name refers to the grammatical number of the attractor (in structures where attractors are present).

		Tense	
		Singular	Plural
Person	First	*run*	*run*
	Second	*run*	*run*
	Third	***runs***	*run*

Table 4.1: English present-tense forms of the verb *to run*.

of a verb takes on an unmarked present tense form of the verb, while the third-person singular takes on a marked form. This is illustrated in Table 4.1 for the verb *to run*. Since the unmarked form of the verb is often much more frequent, it makes sense that the language model may learn a preference for unmarked verbs during training time. This asymmetrical treatment of singular and plural agreement is discussed further at the end of this section.

Grammaticality increases when attractors and subjects match. We expect to see an improvement in grammaticality when subjects and attractors match as there is no potentially conflicting evidence for the number of the verb. Although this effect is dominated by the disparity in plural and singular subject grammatically, the disparity in grammaticality is significantly less when the attractor is singular. This effect is more apparent when considering the relative sizes of bars of a single color in Figure 4.2. For plural subjects, plural attractors result in higher grammaticality, and similarly for singular subjects, singular attractors result in higher grammaticality. The same effect is seen both in across and withing RC cases, further supporting our hypothesis.

Interposing attractors decrease grammaticality, regardless of plurality. However, note that this is not true of distractors: placing adverbs between the subject and verb tends to have little effect, even though the 'across two distractors' structure places the same token distance between subject and verb as 'across a PP'. This means that distance between subject and verb is less important than the type of the structure separating them.

Subject-preceding attractors have subject-dependent impacts on grammaticality. In the 'within singular RC' structure, grammaticality is reduced for both singular and plural subjects compared to the 'simple agreement' structure. However, 'within plural RC' has a polarizing effect:

in the singular subject case, grammaticality is significantly reduced when subject and attractor do not match in number, but in the plural subject case, grammaticality is instead greatly increased when the subject and attractor do match.

This finding lends support to the hypothesis from that GPT-2 treats singular and plural agreement asymmetrically. Jumelet et al. [2019] report related findings that LSTM LM's have "default" verb forms for number agreement. In particular, they find that the singular verb form serves as the default, and additional evidence, such as a plural subject, is necessary to produce a plural verb prediction. This same type of asymmetrical and perhaps "default" preference for a verb form in Transformers may contribute to the patterns we see in our experiments.

4.3 Natural Indirect Effect

Having investigated the total effects of changing the prompt subject number, we now turn to *indirect effect* in order to investigate the contributions of individual model components as mediators of syntactic number agreement. Given an independent variable X, a dependent variable Y, and a mediating variable Z, we want to determine the indirect effect of Z on Y. Suppose that z is the value that Z takes when $X = x^*$, y_{x^*} is the value Y takes when $X = x*$, and y_x is the value Y takes when $X = x$. We can intervene on Z, setting it to z and then find the indirect effect by calculating

$$\frac{y_{xz} - y_{x^*}}{y_{x^*}} \tag{4.5}$$

where y_{xz} is the value Y takes when $X = x$ and $Z = z$.

In our experiments, we are interested in determining which model components mediate number agreement by measuring the indirect effects of model compnents on the output of the model. We can find this using the above formulation by allowing Y to be the output metric defined in Equation 4.1, setting X to the interventions `swap-number` and `null`, and setting Z to be the output of a component within the model.

To formalize this, the natural indirect effect (NIE) is the change in effect when a prompt u is unchanged, but a model component $\mathbf{z}$ (e.g., a neuron) is

set to the value it *would have taken* if the intervention had occurred.

$$\overline{\text{NIE}}(\texttt{swap-number, null}; y, \mathbf{z}) = \mathbb{E}_{u,v}\left[\frac{y_{\texttt{null},\mathbf{z}_{\texttt{swap-number}}(u,v)}(u,v)}{y_{\texttt{null}}(u,v)} - 1\right] \tag{4.6}$$

This allows us to evaluate the contribution of specific parts or regions of a model to the syntactic preferences we observe. More specifically, we can measure to what extent the total effect of swapping the subject on inflection preferences can be attributed to specific neurons. Here, we independently analyze the individual neuron NIEs (though future work could also investigate intervening on sets of neurons simultaneously). We also attempt to analyze attention heads, though we find that they do not present consistent interpretable results. We take a closer look at attention heads in section 4.4.

Based on the findings of prior probing work on dependency parsing [Hewitt and Manning, 2019], we hypothesize that NIEs will peak in the upper-middle layers for all models. Additionally, since GPT-2 and Transformer-XL always process input left-to-right, we expect differing indirect effect results between different structures, but similar indirect effect results overall across these two model architectures. Because XLNet is exposed to all word order permutations of its input sentences during training, we hypothesize that results may be more similar across syntactic structures, and that it will display different indirect effect trends compared to the other two models.

4.3.1 Results

For each model and structure, we select the 5% of neurons with the highest NIE in each model layer; Figure 4.3 compares NIEs across model sizes, and Figure 4.4 compares NIEs across structures for GPT-2 Medium.[2] When weights are randomized, indirect effects peak at layer 0—albeit at values perhaps too small to be meaningful—and then remain close to 0 in higher layers. This indicates that the vast majority of the indirect effect observed for trained models is due to what is learned, rather than chance.

[2] We also produced figures using *all* neurons. When doing so, the contour of the graph across layers did not change, but the magnitudes were lower since we average over more neurons.

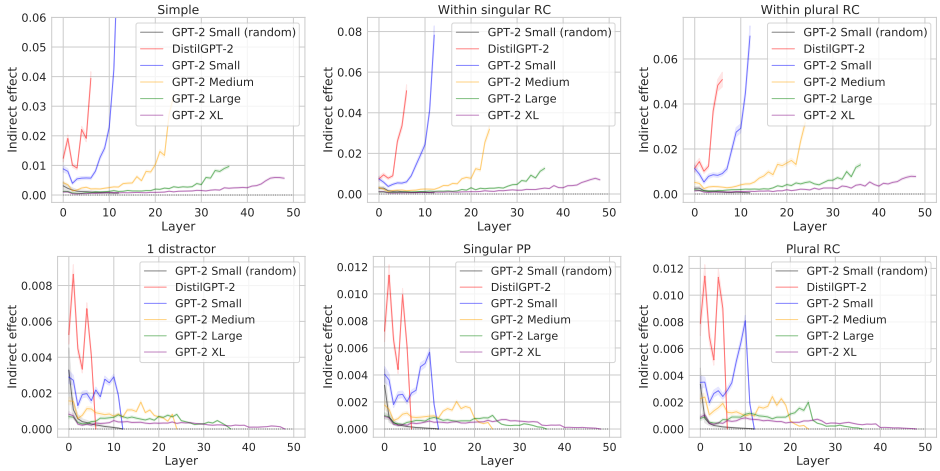

Figure 4.3: Natural indirect effects the top 5% of neurons in each layer of various GPT-2 sizes. Each figure focuses on a single structure and compares across GPT-2 sizes.

GPT-2 and TransformerXL have bifurcated number agreement mechanisms. We observe two distinct layer-wise patterns for these models. In structures where the target verb directly follows the subject ('simple agreement' and 'within RC', the top 3 plots in Figure 4.3), NIEs continually increase in higher layers. Conversely, for structures with subject-verb separation ('across one/two distractor(s)', 'across PP', and 'across RC', the bottom 3 figures in Figure 4.3), NIEs peak at layer 0 and (more notably) in the upper-middle layers. The effect then decreases sharply, reaching 0 in the uppermost layers.

Even a single token between subject and verb brings about this second indirect effect contour, indicating that **distance is a less important factor than the presence of *any* separation in invoking this second syntactic agreement mechanism**. We again see that when weights are randomized, effects peak (at low values) in lower layers, quickly dropping to zero and remaining low in higher layers.

These two distinct indirect effect contours may indicate two distinct subject-verb agreement mechanisms, depending on subject-verb separation. This resembles a finding of Lakretz et al. [2019], who found distinct systems for short- and long-distance agreement in LSTMs.

The peak NIE is lower here than for structures where there is no separation,

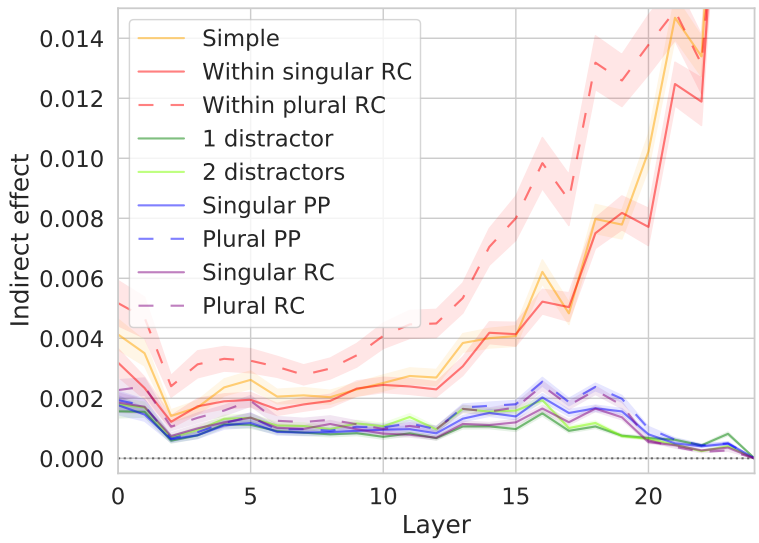

Figure 4.4: Natural indirect effects of the top 5% of neurons in each layer of GPT-2 Medium.

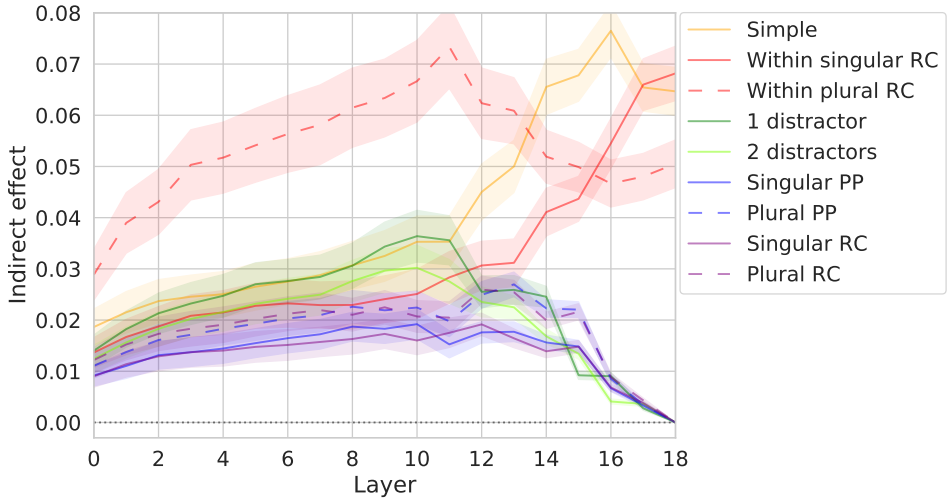

Figure 4.5: Natural indirect effects of the top 5% of neurons in each layer of Transformer-XL.

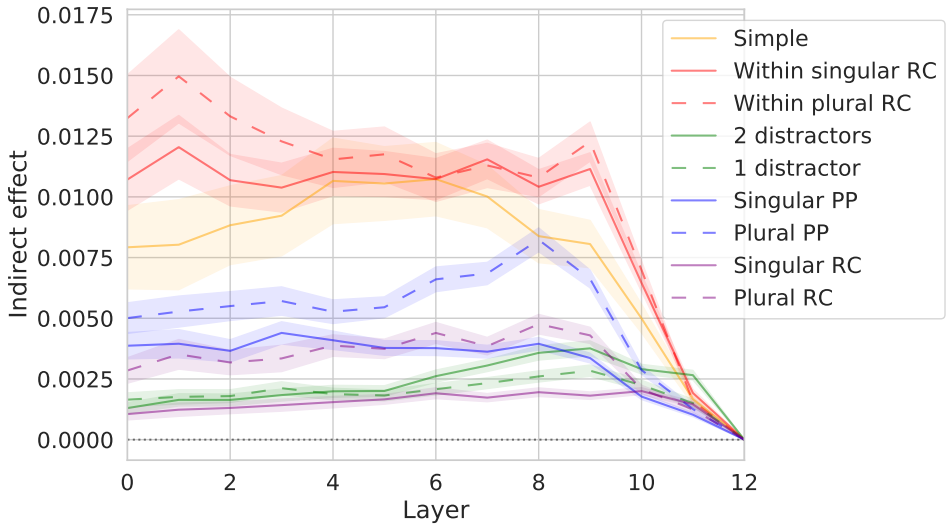

Figure 4.6: Natural indirect effects of the top 5% of neurons in each layer of XLNet.

perhaps indicating that syntactic agreement information is more distributed among neurons when separation occurs. For each structure, the maximum NIE per layer is always lower for larger models. Peaks in NIEs are also more distributed across layers for larger models. **This suggests that structural knowledge is more concentrated in fewer neurons with stronger inflectional preferences in smaller models, whereas it is more distributed across neurons in larger models.** Nonetheless, the overall contour of NIEs is similar across model sizes for a given structure, indicating that the **underlying mechanisms of agreement are similar across model sizes**.

4.3.2 Attention Heads

Trends are generally less clear for attention heads. Each structure appears to have completely different heads with significant indirect effects; even changing the number of the attractor has a large effect on which heads are implicated in observed behavior. See section 4.4.

Perhaps a reason for this is that most of the total effect does not flow through the attention heads. We find that the sum of the indirect effects across heads averaged across prompts is always near-zero (and is often small and negative),

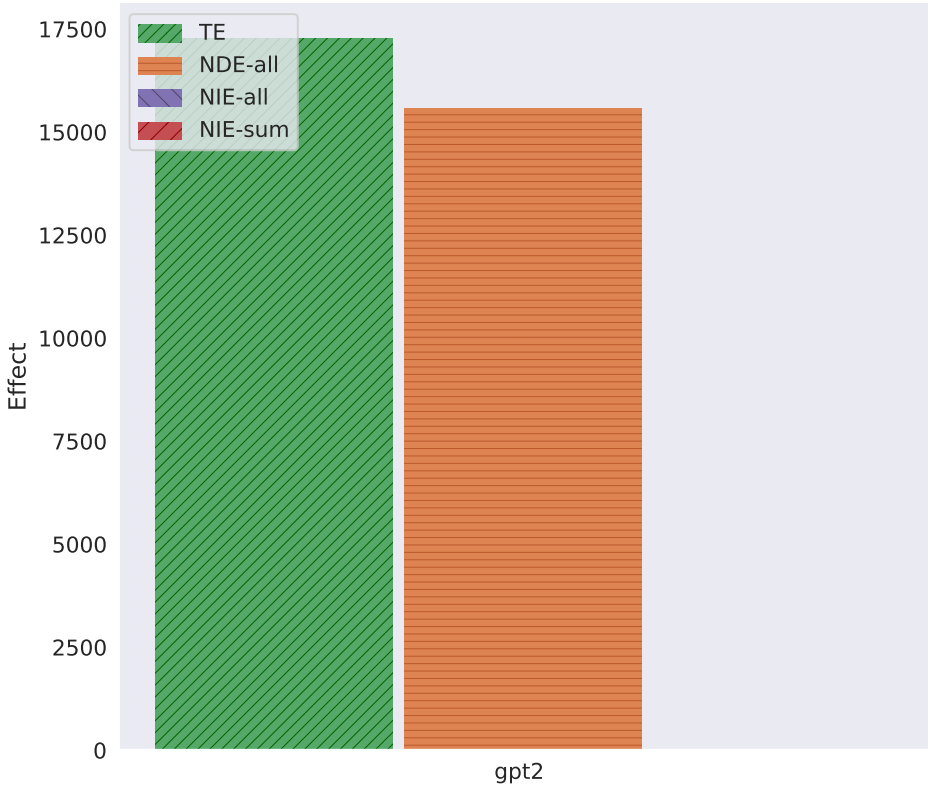

Figure 4.7: Total (TE), direct (NDE-all), and indirect effects (NIE-all and NIE-sum) for attention heads for the 'across a distractor' structure. NIE-all (the indirect effect from simultaneously intervening on all attention heads) and NIE-sum (the indirect effect from summing independent and separate interventions on each attention head) are near zero and thus not visible.

as is the indirect effect attained when intervening on all attention heads simultaneously. Meanwhile, the direct effect across heads is very large and captures most of the total effect. See Figure 4.7. This indicates that most of the total effect does not flow through the attention heads, but through other model components (e.g., word representations, neurons). Thus, perhaps what specific heads are displaying across structures is noise, rather than meaningful effects.

To test whether attention is simply robust to the `swap-number` intervention, we define an alternate intervention `zero`: rather than swapping the number of the subject, we set a specific attention head **z**'s value to 0 and observe how this changes the effect on the unchanged prompt. This may be interpreted as the controlled direct effect, as defined in Pearl [2001].

For this intervention, trends are much more consistent across structures. Head 0-10 is always strongly implicated in behavior; since this is in the bottom layer, this indicates that much of attention's contribution to syntax is lexicalized. In other words, attention's contribution to subject-verb agreement may largely be based on which word embeddings are present in the input, rather than any compositional structural information that would be present in the upper layers.

For all structures except where we have adverbial distractors, we see consistent positive indirect effects in the uppermost layers as well. No single attention head appears strongly implicated (though head 11-0 is always positive), but the layer effect is consistently positive and sometimes nears the magnitude of the layer effect of layer 0. This indicates that more abstract structural information is quite distributed across attention heads in the uppermost layers, but nonetheless may be present.

Comparison to Other Architectures

We also investigate the neuron NIEs of Transformer-XL (Figure 4.5) and XLNet (Figure 4.6) to observe whether syntax is represented in a similar manner across models (for total effects across architectures, see section 4.8).

We observe similar dichotomous agreement mechanisms across GPT-2 and Transformer-XL. 'Simple agreement' and 'within RC' show similar mechanisms of subject-verb agreement, whereas separating subject and verb lead to very different mechanisms of agreement. Moreover, the contour of the NIEs across layers is similar for these two architectures. This

supports our hypothesis that GPT-2 and Transformer-XL encode syntax in a similar manner.

Conversely, indirect effects in XLNet are different to those seen in GPT-2. We do not observe the same dichotomous behavior between subject-verb adjacent and subject-verb non-adjacent structures; rather, the overall contours are all the same. All of the indirect effects approach 0 in the final layer. Overall, this resembles the way GPT-2 and Transformer-XL treat structures where subject and verb are not adjacent. This could be due to XLNet seeing all word order permutations of training sentences, thus making it more robust to the distance between subject and verb. This supports our hypothesis that XLNet encodes syntax in a different manner to GPT-2.

Neuron Overlap Across Structures

To investigate whether the same neurons are implicated in subject-verb agreement across structures, we select the top 5% of neurons per layer by NIE and calculate the proportion of these high-NIE neurons that overlap between each pair of structures.

Does the extent of neuron sharing across structures correlate with human intuitions of syntactic similarity? To address this question, we compute hypothesized syntactic similarities between structures based on the following linguistic features: distance between subject and verb; presence of adverbial distractors, a relative clause, prepositional phrase, and/or a noun attractor; and the number of the noun attractor when present. Appendix 4.7.1 details how we calculate similarity.

To quantify the similarity of the hypothesis matrix and a neuron overlap matrix, we calculate the L_1 norm of the element-wise difference between the lower-left triangle of both matrices, as the matrices are symmetric (using the L_2 norm does not change which layer has the lowest difference norm). We exclude the diagonal.

For each model, we present neuron overlaps for the layer with the lowest difference norm to the hypothesis (Figure 4.8; for an analysis of layer-by-layer overlap change for GPT-2, see subsection 4.7.2). The lowest difference norms are 443 (GPT-2), 510 (Transformer-XL), and 486 (XLNet). GPT-2 Medium's overlap across structures at layer 21 (of 24) is visually similar to the hypothesis, meaning that **this layer in GPT-2 shares neurons for subject-verb agreement across structures in a way that aligns with**

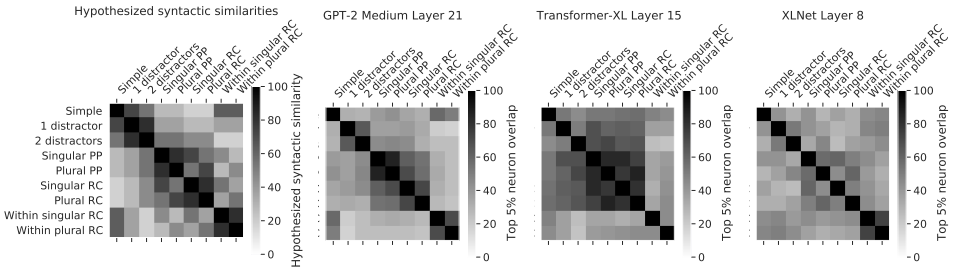

Figure 4.8: Hypothesized syntactic similarity across structures (left), as well as the overlap of the top 5% of neurons per-structure by indirect effect for GPT-2 (center-left), Transformer-XL (center-right), and XLNet (right).

human intuitions about syntactic similarity. Interestingly, it learns to do this without receiving explicit syntactic supervision during training.

Layer 15 (of 18) of Transformer-XL displays similar trends to GPT-2, though the extent of overlap is higher across structures in general here. There is more significant overlap between the adverbial distractor structures and the structures that contain attractors. 'Simple agreement' also has more overlap with structures containing attractors than 'within RC', which is contrary to the hypothesis. Also note how 'across singular RC' has more overlap with 'across PP' than 'across plural RC' (and vice versa for 'across plural RC'), indicating that **the number of the attractor is more salient to Transformer-XL than the structure of the phrase containing the attractor.**

Meanwhile, Layer 8 (of 12) of XLNet appears noisier. There is slightly more overlap between structures across noun attractors, but the extent of overlap is smaller compared to other models, suggesting that there are more neurons specialized to specific structures. However, the indirect effect findings for XLNet suggest a more unified mechanism for syntactic agreement across all structures; if this were the case, we would expect neuron overlap to be high, and for the extent of overlap to be similar across *all* structures, rather than being higher between more similar structures. We observe the latter, but not the former. Regardless, both observations further support our hypothesis that XLNet uses different mechanisms to encode syntax than the other two architectures.

4.4 Attention Head Indirect Effects

Here, we present the mean indirect effects across prompts of each attention head for a sample of structures for GPT-2 Small, under both the `swap-number` intervention (Figure 4.9) and `zero` intervention (Figure 4.10; intervention defined below).

For the `swap-number` intervention, we do not observe any consistent trends across each structure, except that attention heads in the upper-middle layers seem to account for most of the positive *and* negative indirect effects. Head 10-9 seems to have negative indirect effects for most structures where there is a separation between subject and verb, except 'across plural RC'. However, we do not observe strong indirect effects for head 10-9 when subject and verb are adjacent. Head 11-11 has the most consistently positive indirect effects across structures, though its magnitude is typically low.

Indirect effects are largely positive for 'within plural RC', but otherwise, indirect effects are relatively evenly split between positive and negative. Indeed, the sum of indirect effects across heads for most structures is very close to 0, with many sums being a low-magnitude negative number. This indicates that these attention indirect effects may simply be noise.

Because attention heads seem robust to swapping the number of the subject, we also define the `zero` intervention. Here, we do not change u, but set the attention head's value equal to 0 and observe how this changes the effect; this has an interpretation as the *controlled* indirect effect from Pearl [2001]. Here, trends are more consistent across structures, attractor numbers, and types of distractors. Head 0-10 is always strongly implicated in behavior; since this is in the bottom layer, this suggests that attention's contribution to syntax is lexicalized. Qualitative analysis reveals that head 0-10 and head 2-8 always focus on the 2nd and 5th words in the prompt, respectively. Thus, attention's contribution to subject-verb agreement may largely be based on where important tokens appear in the input, rather than any abstract structural information that would be composed in the upper layers.

However, for all structures except where we have adverbial distractors, we see consistent positive indirect effects in the uppermost layers as well. No single attention head appears strongly implicated, but the layer effect is consistently positive and sometimes nears the magnitude of that in the lower layers. This indicates that more abstract structural information is quite distributed across attention heads in the uppermost layers, but nonetheless may be present. Future work should investigate other interventions to better

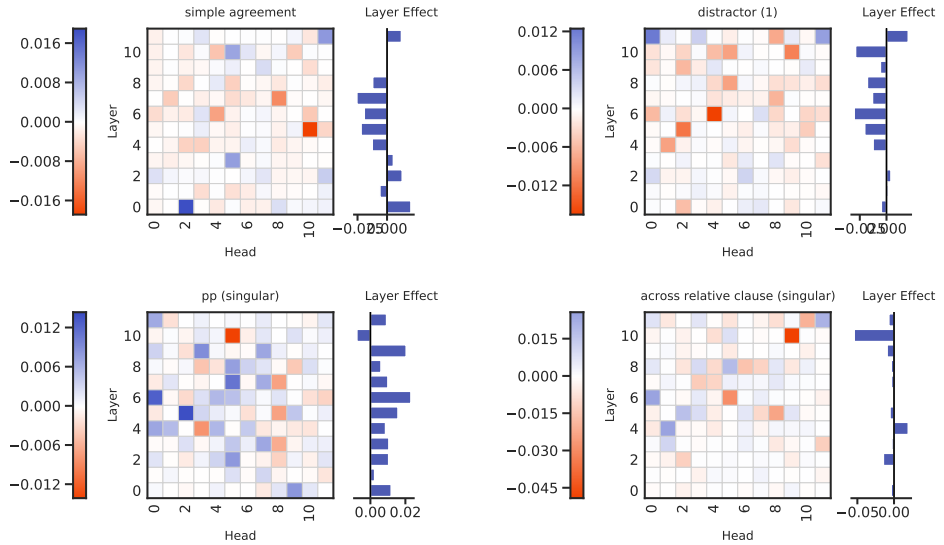

Figure 4.9: Attention indirect effects for GPT-2 Small under the `swap-number` intervention.

understand these mechanisms of agreement.

4.5 Adverbs Increase the Probability of Correct Verbs More Than Incorrect Verbs

Here, we show that separating the subject and verb with adverbs tends to increase the probability of the verb (Figure 4.11). Regardless of whether the subject and verb agree, adding adverbs does always increase the probability of verbs. Note the log scale: we observe visually similar increases in log probabilities after adding 1 or 2 adverb distractors, but visually similar differences at higher points in the graph are actually much larger increases. This supports our hypothesis that adverbs increase all verb probabilities—but also that they increase correct inflection probabilities more.

4.6 The (Non-)Impact of Complementizers

Here, we investigate the effect of including or excluding the complementizer *that* for the 'across RC' and 'within RC' structures, observing both TEs (Figure 4.12) and neuron indirect effects (Figure 4.13). While we expect lower

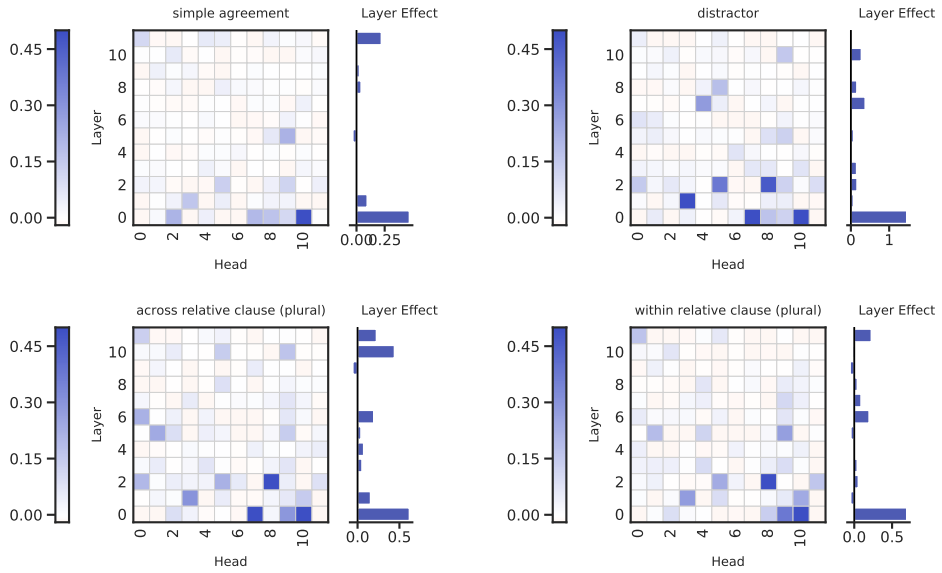

Figure 4.10: Attention indirect effects for GPT-2 Small under the `zero` intervention.

TEs when the complementizer is absent, we observe only minor reductions in total effects in 'across RC'; this holds across model sizes. For 'within RC', however, trends are size-dependent. DistilGPT-2, GPT-2 Small, and GPT-2 Large appear mostly robust to the presence or absence of the complementizer, though GPT-2 Small does have lower total effects in the 'across plural RC' structure when *that* is absent. Meanwhile, GPT-2 Medium more strongly prefers correct inflections when *that* is absent. It is not immediately clear why this is the case, because deleting the complementizer introduces more ambiguity.

One might expect that the exclusion of the complementizer will change neuron behavior when agreeing subjects and verbs. However, there does not appear to be any significant difference in magnitude or contour of the indirect effects across layers when including or excluding it. Thus, while excluding the complementizer can make subject-verb agreement slightly more difficult for LMs [Marvin and Linzen, 2018], it does not appear to change the mechanisms through which subject-verb agreement happens in the model.

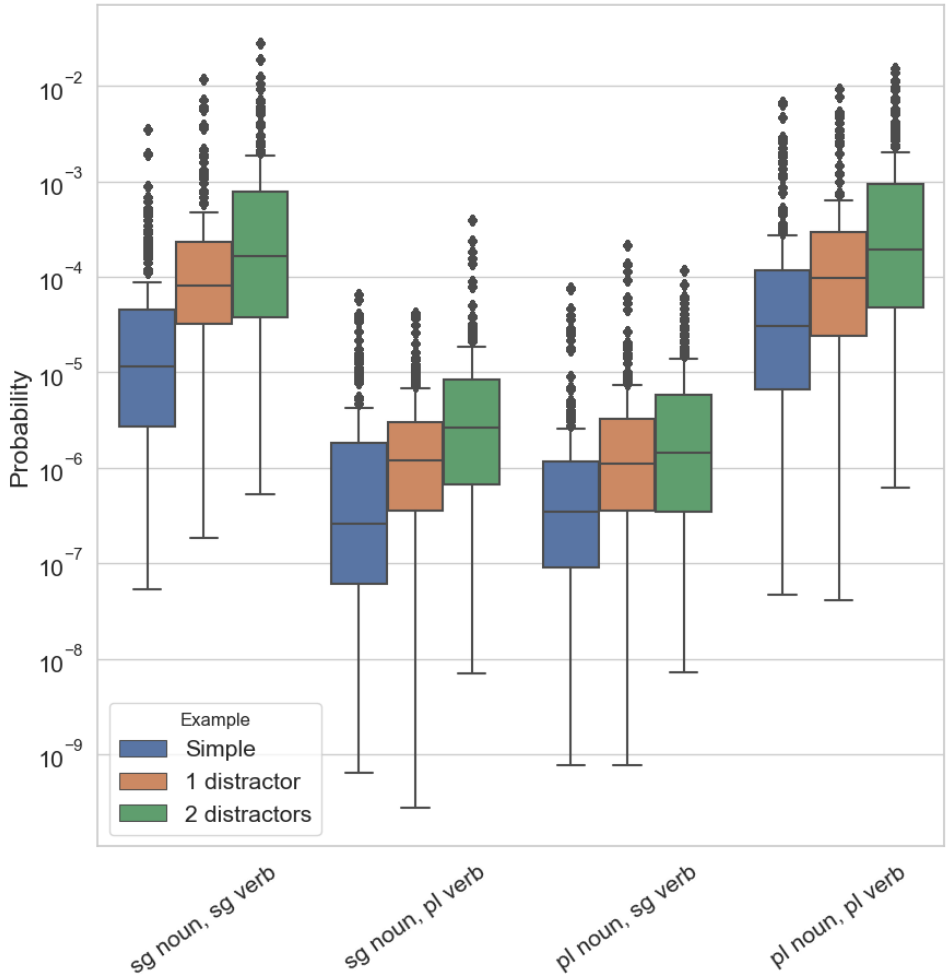

Figure 4.11: The distribution of target verb probabilities for 'simple agreement', 'across one distractor', and 'across two distractors'. Note that the y-axis uses a log scale.

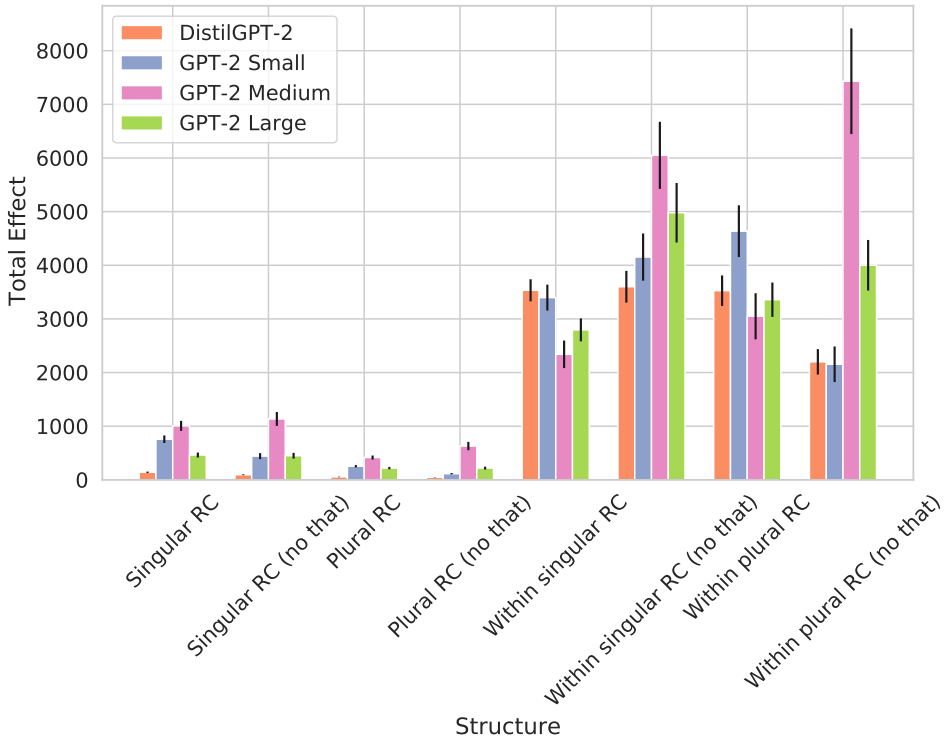

Figure 4.12: Total effects for 'across relative clause' and 'within relative clause' structures, with and without the complementizer *that*.

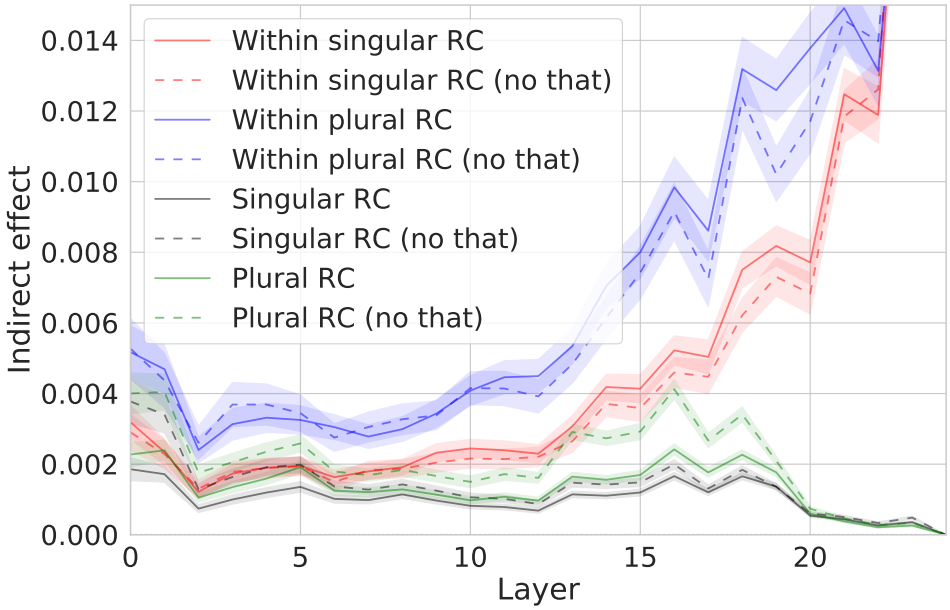

Figure 4.13: Indirect effects for 'across relative clause' and 'within relative clause' structures, with (solid lines) and without (dashed lines) the complementizer *that*.

4.7 Additional Neuron Overlap Details

4.7.1 Hypothesizing Syntactic Similarity

To generate the hypothesis similarity matrix between structures, we choose a set of features given in Table 4.2 that capture important syntactic information according to our intuition. Most of the features are binary; however, we also include a ternary feature and a numerical feature. The ternary feature, "attractor number", can take on values SG, PL, and 0 when there is no attractor.

To compute the similarity of two structures, we first sum the difference for each feature. For the binary and ternary features, the difference is 0 if the features have the same value, otherwise 1. For the numerical feature, we take the absolute value of the difference between distances, scaled to a value between 0 and 2. We scale the similarity to reduce the impact of the numerical feature on the total similarity.[3] Finally, we take the maximum possible difference across all pairs of structures, and subtract each pairwise distance from the maximum to obtain similarity scores. We normalize the similarities to the range [0, 100] by dividing similarities by the maximum possible similarity score; this is to make them more comparable to the neuron overlap matrices.

4.7.2 Neuron Overlap Across Layers

Here, we present neuron overlaps across all layers of DistilGPT-2 (Figure 4.14). We first note that the overall extent of neuron overlap across structures tends to increase up to the upper-middle layers, before sharply decreasing in the highest layer to near-zero values. We find that this trend holds for all other sizes of GPT-2, as well as Transformer-XL; generally, overlaps continue to increase until the upper-middle layers, decreasing slightly in the second-highest layer and decreasing sharply to zero in the highest layer. Layer-by-layer difference (L_1) norms are presented in Table 4.3.

4.8 Total Effects Across Architectures

Figure 4.15 presents total effects for all structures across architectures. The magnitude of total effect is generally similar for XLNet and GPT-2 (except

[3] We initially scaled this to be within the range [0, 1] like the other features, but this caused "within relative clause" to have high similarity to "across PP" and "across a relative clause". Thus, we increase its impact for more human-like hypotheses.

Feature	Type
Subject and verb separated	binary
Tokens between subject, verb	numerical
Has adverbial distractor(s)	binary
Has noun attractor	binary
Attractor number	ternary
Has relative clause	binary
Has prepositional phrase	binary

Table 4.2: Features (and their types) used in calculating hypothesized syntactic similarity.

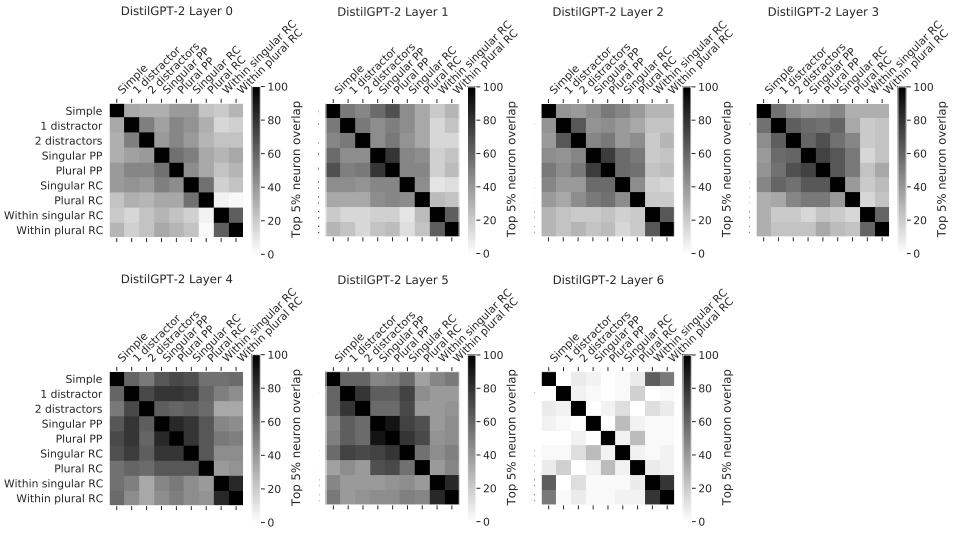

Figure 4.14: Overlap between structures of the top 5% of neurons in each of layer of DistilGPT-2 by indirect effect.

Layer No.	Diff. Norm	Layer No.	Diff. Norm
0	677	4	627
1	652	5	510
2	565	6	1301
3	583		

Table 4.3: Difference L_1 norms between the hypothesis matrix and each layer of DistilGPT-2.

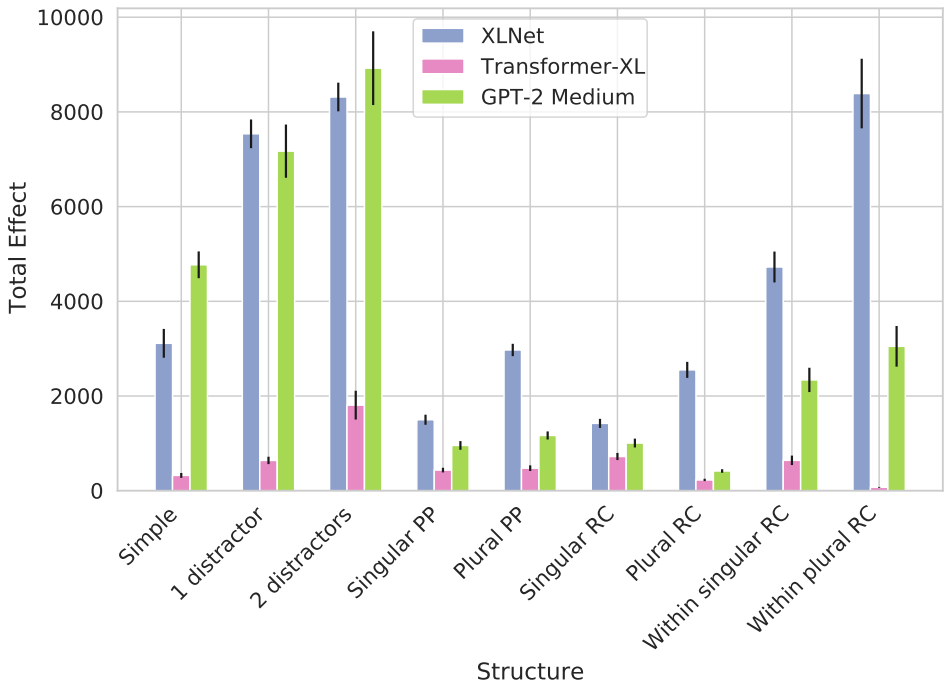

Figure 4.15: Total effects across structures by architecture (ordered by increasing number of layers/increasing parameterization).

when dealing with relative clauses), whereas total effects are much smaller for Transformer-XL. It seems that parametrization and model depth do not correlate well with total effects.

Perhaps the effects for Transformer-XL are smaller due to the longer effective contexts it has, which could make it prone to assigning smaller probabilities to a larger set of tokens than GPT-2 while still behaviorally performing well. However, it is harder to explain the similarity between GPT-2 and XLNet, given the great differences between them in training and the divergence in their behavior as revealed by the indirect effect results in §4.3.2.

Chapter 5

Conclusion

This study applied causal mediation analysis to discover and interpret the mechanisms behind syntactic agreement in pre-trained neural language models. Our results reveal the location and importance of various neurons within various models, and provide insights into the inner workings of these LMs.

Future work should investigate other interventions to better understand how other model components (e.g., attention weights/heads) give rise to subject-verb agreement. More broadly, it would be beneficial to investigate other syntactic phenomena such as filler-gap dependencies and negative polarity items.

Bibliography

Y. Belinkov and J. Glass. Analysis methods in neural language processing: A survey. *Transactions of the Association for Computational Linguistics*, 7: 49–72, Mar. 2019. doi: 10.1162/tacl_a_00254. URL https://www.aclweb.org/anthology/Q19-1004.

E. A. Chi, J. Hewitt, and C. D. Manning. Finding universal grammatical relations in multilingual BERT. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5564–5577, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.493. URL https://www.aclweb.org/anthology/2020.acl-main.493.

Z. Dai, Z. Yang, Y. Yang, J. Carbonell, Q. V. Le, and R. Salakhutdinov. Transformer-XL: Attentive language models beyond a fixed-length context, 2019.

J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019a. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL https://www.aclweb.org/anthology/N19-1423.

J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019b. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL https://www.aclweb.org/anthology/N19-1423.

A. Geiger, K. Richardson, and C. Potts. Neural natural language inference models partially embed theories of lexical entailment and negation. In *Proceedings of the Third BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 163–173, 2020.

M. Giulianelli, J. Harding, F. Mohnert, D. Hupkes, and W. Zuidema. Under the hood: Using diagnostic classifiers to investigate and improve how language models track agreement information. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 240–248, 2018.

Y. Goldberg. Assessing BERT's syntactic abilities. arXiv preprint 1901.05287, 2019. URL http://arxiv.org/abs/1901.05287.

K. Gulordava, P. Bojanowski, E. Grave, T. Linzen, and M. Baroni. Colorless green recurrent networks dream hierarchically. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1195–1205, New Orleans, Louisiana, June 2018. Association for Computational Linguistics. doi: 10.18653/v1/N18-1108. URL https://www.aclweb.org/anthology/N18-1108.

J. Hewitt and P. Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743, Hong Kong, China, Nov. 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1275. URL https://www.aclweb.org/anthology/D19-1275.

J. Hewitt and C. D. Manning. A structural probe for finding syntax in word representations. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4129–4138, 2019.

S. A. B. Hill. The environment and disease: Association or causation? *Proceedings of the Royal Society of Medicine*, 58(5):295–300, 1965. doi: 10.1177/003591576505800503. URL https://doi.org/10.1177/003591576505800503. PMID: 14283879.

J. Hu, J. Gauthier, P. Qian, E. Wilcox, and R. Levy. A systematic assessment of syntactic generalization in neural language models. In *Pro-*

ceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pages 1725–1744, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.158. URL https://www.aclweb.org/anthology/2020.acl-main.158.

J. Jumelet, W. Zuidema, and D. Hupkes. Analysing neural language models: Contextual decomposition reveals default reasoning in number and gender assignment. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 1–11, Hong Kong, China, Nov. 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1001. URL https://www.aclweb.org/anthology/K19-1001.

Y. Lakretz, G. Kruszewski, T. Desbordes, D. Hupkes, S. Dehaene, and M. Baroni. The emergence of number and syntax units in LSTM language models. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 11–20, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1002. URL https://www.aclweb.org/anthology/N19-1002.

T. Linzen, E. Dupoux, and Y. Goldberg. Assessing the ability of LSTMs to learn syntax-sensitive dependencies. *Transactions of the Association for Computational Linguistics)*, 4:521–535, 2016. URL https://transacl.org/ojs/index.php/tacl/article/view/972.

J. Mahoney. After kkv: The new methodology of qualitative research. *World Politics*, 62(1):120–147, 2010. ISSN 00438871, 10863338. URL http://www.jstor.org/stable/40646193.

R. Marvin and T. Linzen. Targeted syntactic evaluation of language models. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1192–1202, Brussels, Belgium, Oct.-Nov. 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1151. URL https://www.aclweb.org/anthology/D18-1151.

M. Maziarz. *The Philosophy of Causality in Economics: Causal Inferences and Policy Proposals*. Routledge INEM Advances in Economic Methodology. Taylor & Francis, 2020. ISBN 9781000069105. URL https://books.google.com/books?id=NGIPEAAAQBAJ.

A. Mueller, G. Nicolai, P. Petrou-Zeniou, N. Talmina, and T. Linzen. Cross-linguistic syntactic evaluation of word prediction models. In *Pro-*

ceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pages 5523–5539, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.490. URL https://www.aclweb.org/anthology/2020.acl-main.490.

J. Pearl. Direct and indirect effects. In *Proceedings of the Seventeenth conference on Uncertainty in artificial intelligence*, pages 411–420, 2001.

A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners. 2019. URL https://github.com/openai/gpt-2.

J. M. Robins. Semantics of causal dag models and the identification of direct and indirect effects. *Oxford Statistical Science Series*, pages 70–82, 2003.

J. M. Robins and S. Greenland. Identifiability and exchangeability for direct and indirect effects. *Epidemiology*, pages 143–155, 1992.

V. Sanh, L. Debut, J. Chaumond, and T. Wolf. DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter, 2020.

M. Sundermeyer, R. Schlüter, and H. Ney. LSTM neural networks for language modeling. In *Thirteenth annual conference of the international speech communication association*, 2012.

M. van Schijndel, A. Mueller, and T. Linzen. Quantity doesn't buy quality syntax with neural language models. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5835–5841, Hong Kong, China, Nov. 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1592. URL https://www.aclweb.org/anthology/D19-1592.

A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in neural information processing systems*, pages 5998–6008, 2017.

J. Vig, S. Gehrmann, Y. Belinkov, S. Qian, D. Nevo, Y. Singer, and S. Shieber. Causal mediation analysis for interpreting neural NLP: The case of gender bias, 2020.

E. Wilcox, R. Levy, T. Morita, and R. Futrell. What do RNN language models learn about filler–gap dependencies? In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 211–221, Brussels, Belgium, Nov. 2018. Associa-

tion for Computational Linguistics. doi: 10.18653/v1/W18-5423. URL https://www.aclweb.org/anthology/W18-5423.

E. Wilcox, P. Qian, R. Futrell, M. Ballesteros, and R. Levy. Structural supervision improves learning of non-local grammatical dependencies. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3302–3312, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1334. URL https://www.aclweb.org/anthology/N19-1334.

Z. Yang, Z. Dai, Y. Yang, J. Carbonell, R. Salakhutdinov, and Q. V. Le. XLNet: Generalized autoregressive pretraining for language understanding, 2020.